

- **10th October**
 - Started working on the assignment from scratch, initially following the broad outline (3 smaller tasks: single-threaded, multithreaded & parsing HTTP reqs)
 - I've got something that I think should work (for single-threaded), but the server doesn't seem to be accepting connections
 - The problem was that I was passing `arg[0]` not `arg[1]` as the port, which (whilst it worked), means that the socket was bound to a **completely** different port than expected because I was using *stoi* rather than *strtoul*...
- **11th October**
 - Added some multithreading, which seems to be working. When a client disconnects, it doesn't work properly though.
 - I was closing `clisockfd` rather than `*clisockfd`, but I'm passing a pointer to the file descriptor to the `handleConnection` method - Now I've fixed this.
 - I'm going to focus on the OS assignment and my FYP for a bit and then come back to this - I feel like I'm making good progress.
- **17th October**
 - Added structs and method stubs that will be used in methods to handle everything to do with HTTP responses - my plan is to initially respond with a successful request (containing the index page) regardless of the incoming request so I can test that it would successfully be displayed.
- **18th October**
 - I've added the functionality to timestamp requests, and have implemented most of the methods required for creating/deleting/constructing/sending an HTTP response
 - The server now correctly serves a 200 response with the index page (stored as a `char *` in the program), which is displayed by the browser. I've ran into a couple of memory errors that I need to diagnose though...
 - The memory errors were caused because `strlen` doesn't count the null terminator of the string in the returned length. I've now fixed this.
 - I've added the appropriate method stubs and structs for handling HTTP requests - I plan to start handling GET requests soon, after I implement the *parse_http_request* method.
 - Whilst working, I noticed a potentially insecure piece of code where I'd erroneously used `strcpy` rather than `strncpy` on the body of the request that's provided by the user. This could have been used for a buffer overflow attack.
- **24th October**
 - I've started work on actually handling HTTP requests when they're sent to the server
 - Namely, I've started implementing the *parse_http_request* method. I was originally going to use `strtok(2)` for this, but because it stores state, it isn't threadsafe. Instead, I'm using `strtok_r(3)`, which is threadsafe.
 - For now, I'm just going to parse the request line of the request - I can get enough information from this to fulfil at least GET requests.

- **25th October**

- I've written a `ReadFile` method which reads a given file on the system into a `char *` - I'm going to use this to serve files from the server.
- Because I want to restrict access to files requested, I've created a simple whitelist system and *servable* method, that just checks through a list (`files.txt`) of whitelisted files.
- I'm now correctly handling 404 requests too - If a file isn't servable (i.e. it isn't whitelisted), the server will send back a 404 request.
- If the `http_verb` isn't recognised, the server now sends back a 400 error.
- I've tried reading in image files, but it doesn't seem to read or send the whole file, just the header...

- **26th October**

- Continuing work on the issues with images - the problem seems to be down to the fact that images can contain null bytes that aren't terminating
- The first problem seems to be that I'm using `strlen` to get the length of the contents of the file once it has been read into a `char *` - I've changed how I'm doing this now. This still doesn't seem to have fixed the overarching issue though
- Another problem is that I've been using *strncpy* in the *add_http_response_body* method, but it only copies up to **either** the specified length, or when a null character is reached. I've switched to using *memcpy* here instead.
- The last piece to the puzzle is that I was writing to the socket using *dprintf* (to write to the fd), and supplying the body of the response as part of the format string - but using the `%s` token. This meant that only the characters up to the first null character were ever sent. I've changed this around to write everything except the body using *dprintf*, and then doing a second *write* with the body.
- Extended the server to handle (most) other MIME types.
- I've also added a system to keep track of allocated memory so that in the case of early termination etc. it's all cleaned up prior to exit.
- Compiled everything with `pedantic` and made sure it's all compiling properly with the Makefile I've provided.

- **27th October**

- I've started doing some more extensive testing of the server itself - there are a few cases that caused it to crash (such as trying to access a non-servable (or non-existent) resource) - I've now sorted these out and the server seems relatively robust now.
- If I get chance, my next aim is to implement HEAD requests as these aren't too different to GET requests.